

Astera Labs (ALAB) Q2 2026 Earnings Call Transcript

2026-08-04

Astera Labs (ALAB) Q2 2026 Earnings Call Transcript

Source: The Motley Fool full transcript (call held 2026-08-04, published 2026-08-11)

Company: Astera Labs, Inc. (NASDAQ: ALAB) - AI server connectivity chips (PCIe/CXL retimers, AEC cables, AI Fabric switches) **Highlights:** Q2 2026 revenue \$392M (+104% YoY); Scorpio X AI Fabric switches in volume production; Q3 guidance \$550M (+40% QoQ)

Operator: Good afternoon. My name is Holly, and I will be your conference operator today. At this time, I would like to welcome everyone to the Astera Labs Q2 2026 earnings conference call. All lines have been placed on mute to prevent any background noise. After management remarks, there will be a question and answer session. If you would like to ask a question during this time, simply press star followed by the number one on your telephone keypad. If you would like to withdraw your question, press star one again. Thank you. I will now turn the call over to Leslie Green, investor relations for Astera Labs. Leslie, you may begin.

Leslie Green: Thank you, Holly, and good afternoon, everyone, and welcome to the Astera Labs second quarter 2026 earnings conference call. Joining us on the call today is Jitendra Mohan, Chief Executive Officer and Co-founder; Sanjay Gajendra, President and Chief Operating Officer and Co-founder; and Desmond Lynch, Chief Financial Officer. Before we get started, I would like to remind everyone that certain comments made in this call today may include forward-looking statements regarding, among other things, expected future financial results, strategies and plans, future operations, and the markets in which we operate. These forward-looking statements reflect management's current beliefs, expectations, and assumptions about future events, which are inherently subject to risks and uncertainties that are discussed in detail in today's earnings release and the periodic reports and filings we file from time to time with the SEC, including the risks set forth in our most recent annual report on Form 10-K. It is not possible for the company's management to predict all risks and uncertainties that could have an impact on these forward-looking statements or the extent to which any factor or combination of factors may cause actual results to differ materially from those contained in any forward-looking statement. In light of these risks and uncertainties and assumptions, the results, events, or circumstances reflected in the forward-looking statements discussed during this call may not occur, and actual results could differ materially from those anticipated or implied. All of our statements are based on information available to management as of today, and the company undertakes no obligation to update such statements after the call, except as required by law. Also, during this call, we will refer to certain non-GAAP financial measures, which we consider to be important measures of the company's performance. For example, the overview of our Q2 financial results and Q3 financial guidance refer to various non-GAAP financial measures. These non-GAAP financial measures are provided in addition to, and not as a substitute for, financial results prepared in accordance with US GAAP. A discussion of why we use non-GAAP financial measures and reconciliations between our GAAP and non-GAAP financial measures and outlook are available in the earnings release we issued today, which can be accessed through the investor relations portion of our website. With that, I would like to turn the call over to Jitendra Mohan, CEO of Astera Labs. Jitendra?

Jitendra Mohan: Thank you, Leslie. Good afternoon, everyone. thanks for joining our second quarter conference call for fiscal year 2026. Today, I'll provide an update on AI infrastructure market trends, our Q2 results, and recent announcements. I'll turn the call over to Sanjay to discuss Astera Labs' long-term growth profile. Des will cover our Q2 financial results and Q3 guidance. Since our last earnings call, investment in AI infrastructure

has continued to accelerate. Multiple industry forecasts now point to hyperscaler AI and cloud infrastructure spending in 2027 reaching into the trillion-dollar range, with sovereign AI initiatives, inference workloads, and enterprise adoption expanding globally. These secular trends are increasing the need for intelligent connectivity and expanding our long-term growth opportunity. Astera Labs delivered outstanding results in Q2, with record revenue of \$392.4 million, up 27% sequentially and up 104% year-over-year. This strong performance was driven by broad-based strength across our entire product portfolio, reflecting the diversification of our business as we continue to win new designs across multiple customers and product categories. PCIe 6 momentum further accelerated in Q2, representing more than 50% of our total revenue, driven by both Scorpio AI fabric switches and Aries retimers. This milestone demonstrates the maturity and scale of our Gen 6 portfolio and clear market leadership. On the AI fabric front, Scorpio delivered significant growth in Q2. High-radix Scorpio X-Series has entered volume production and will continue to scale materially in the second half of this year. We expect Scorpio to become our largest product family in Q3, which is one quarter ahead of our prior expectations, marking an important strategic milestone. We are now shipping multiple configurations of Scorpio X-Series into scale-up applications to our initial customers, we remain on track to ship Scorpio X-Series to additional customers by year-end. Scorpio P-Series also continued to expand across multiple customers, with new shipments to hyperscalers and AI infrastructure providers. These programs are currently in pre-production and are expected to ramp more meaningfully in 2027. We are also closely engaged with our customers to support their future system architectures with purpose-built AI fabrics, including next-generation PCI Express and UALink protocols. Upcoming AI systems will employ more complex switching topologies to enable larger XPU cluster sizes and higher bandwidth. These designs are expected to drive a material increase in our silicon dollar content per XPU. Our Scorpio fabric switch family is well-positioned for a multi-year growth cycle driven by the following factors. First, higher dollar content per XPU due to greater attach rate and increased capabilities. Second, expanding customer base with exceptional traction and design in pipeline. Third, a greenfield TAM growing rapidly with broadening use cases across training and inferencing. Our signal conditioning business, Aries, delivered record quarterly revenue in Q2, driven by strong adoption of retimers and gearboxes across scale-up and scale-out topologies in AI and general-purpose compute platforms. Aries growth is being fueled by increased AI system deployments, expanding AI accelerator cluster sizes, greater unit attach rates, and higher ASPs associated with the transition to PCIe 6. Taurus also provided strong revenue growth in Q2, with increased units shipped across a variety of AI and general-purpose platforms. During the quarter, we also shipped pre-production volume of our 100G per lane Taurus smart cable modules for 800G AECs for scale-out applications. We expect Taurus to continue its growth trajectory in the second half of the year as 800G deployments expand, driving improved attach rates and higher ASPs compared to 400G solutions. We also expanded our Taurus portfolio with a new family of 3.2T smart retimers and smart redrivers supporting 200G per lane for next-generation Ethernet and UALink connectivity. These are the industry's first solutions to enable smart swap, giving platform designers the flexibility to choose between low-power smart redrivers and long-reach smart retimers without requiring a board redesign. This architectural flexibility is critical as AI systems evolve rapidly and designers need to address increasing signal integrity challenges at higher data rates while optimizing system power consumption. We expect our new high-speed 200G per lane retimers and redrivers to double our Taurus portfolio market opportunity to over \$4 billion by 2030. In CXL, we are seeing renewed momentum as industry is recognizing its potential to unlock performance and utilization benefits in agentic AI applications. For a growing set of memory-intensive inference and agentic AI workloads, memory capacity and utilization are increasingly becoming system constraints, and CXL-attached memory can offer attractive cost performance relative to local HBM. As our hyperscaler customers look to address current memory supply and pricing dynamics, we are also seeing additional opportunities for CXL in mainstream general-purpose server platforms. During Q2, we closed a new design win with our standard Leo memory controller at a U.S. hyperscaler. Looking into 2027, we expect to ship both standard and custom Leo CXL memory controllers in volume across general-purpose compute and AI inferencing applications to two U.S. hyperscalers. Activity and engagement within the CXL ecosystem continues to be robust, and we expect to convert additional designs in the coming quarters. In closing, I'm proud of the results we have delivered, and what excites me most is the trajectory ahead for Astera Labs. Scorpio X-Series is in volume production, and we expect our Scorpio family to become our largest product category by revenue in Q3, a significant milestone in our evolution as a company. At the same time, new and existing designs across Aries, Taurus, and Leo are contributing to our broad-based growth. Intelligent connectivity is now fundamental to AI infrastructure, and we are winning with the customers and platforms that define the market. With proven execution, deep customer partnerships, and growing design win visibility, we are confident in our ability to continue outgrowing the market. With that, let me turn the call over to our President and CEO, Sanjay Gajendra, to provide more details on our

product momentum and strategic initiatives.

Sanjay Gajendra: Thanks, Jitendra, and good afternoon. I'll walk through our product execution and the expanding market dynamics that are positioning Astera Labs to capture a significantly larger share of AI infrastructure spend over the coming years. Our laser focus on execution and innovation has delivered a purpose-built, intelligent connectivity platform to our hyperscaler customers and is solving sophisticated technology challenges against aggressive timelines. This results-oriented mindset has helped us strengthen the trust that our customers and partners place in us. As a result, we continue to unlock many new opportunities in next-generation designs across multiple customers to rapidly expand our top-line growth. Let me deep dive into some of the key product categories, starting with AI fabrics. Our broad portfolio of Scorpio switching solutions is perfectly suited to address the increasing complexity and bandwidth requirements in AI architectures. With leading hyperscalers planning large-scale training and inference deployments with ever-growing accelerator cluster sizes, our Scorpio X-series high-radix fabric switch family provides intelligent AI fabric with industry-leading features to maximize token economics. These capabilities are driving exceptional demand for our Scorpio X smart fabric switches to address the \$20 billion merchant scale-up switching TAM. Scorpio X-series differentiation comes from its intelligence layer. Hardware-accelerated Hyper-cast and in-network compute capabilities deliver up to 2x improvement in collective performance for large-scale training and inference workloads. These capabilities are implemented through our COSMOS software platform, which now extends beyond basic fleet management functionality to enable dynamic traffic shaping and advanced real-time performance management. We're not stopping here. Drawing on our multi-generational partnerships with customers and the lessons learned from deployments at scale, we are building an ambitious roadmap for our scale-up smart fabric switches. Our Scorpio roadmap is closely aligned with our customers' next-generation XPU architectures, enabling more complex topologies, larger cluster sizes, and higher bandwidth. Looking ahead, we expect the content opportunity for Scorpio X series solutions alone to grow well beyond \$1,000 per XPU in future generations of AI platforms. Scorpio P-Series momentum also remains very strong, with both early-stage and production deployments expanding across multiple AI platforms. We are shipping Scorpio P-Series to several customers to enable high-speed connectivity across AI networking applications. Over the long term, we aim to proliferate Scorpio P-Series across a broader set of applications spanning modular AI servers, disaggregated KV cache appliances, storage platforms, and enterprise-level servers. On signal conditioning, Aries delivered record quarterly performance in Q2 and is well-positioned for sustained growth over the long term. PCIe 6.0 retimers and gearboxes have become critical infrastructure for AI servers, where signal integrity directly determines whether systems achieve rated performance. Aries established itself as the gold standard across major XPU and CPU platforms for PCIe 5.0, and this position has continued for PCIe 6.0. We continue to invest in our customers' technology roadmaps, and we expect to have PCIe 7.0 solutions ready to meet their timelines. Taurus is also expected to deliver strong long-term growth, propelled by 800G deployments for AI scale-out networks in the near term and 1.6T and 3.2T-based systems over the long term. We are pleased to report that Taurus has entered pre-production ramp for 100 gig per lane in Q2, which will support platforms utilizing 800G links. Our Taurus portfolio expansion supporting 200G per lane for 1.6T and beyond positions our roadmap through the next Ethernet generation and the forthcoming ramp of UALink, driving continued content gains and above-industry growth for the Taurus product family. Moving to UALink. The consortium's 2.0 specification validates our architectural approach. UALink 2.0 codifies in-network compute, advanced RAS features, and 200 gig SerDes as next-generation scale-up fabric requirements. These directly align with our investment areas. We remain engaged with lead customers and ecosystem partners to enable UALink deployments supported by our upcoming fabric and signal conditioning solutions. We are on track to deliver UALink-enabled Scorpio X Series scale-up switches in 2027. Optical connectivity represents a substantial expansion vector with the potential to unlock tens of billions of dollars of additional market opportunity beyond our current copper connectivity TAM. Over the past two years, we have assembled capabilities spanning analog and mixed signal design, DSP, electronic and photonic IC integration, and optical packaging through organic investment and the aiXscale acquisition. This enables us to deliver complete optical engines in addition to optical components. Our optical strategy unfolds in three phases. First, high-density fiber-attached solutions leveraging our aiXscale core technology enable up to 50-meter fiber connectivity with ultra-low latency. Our initial opportunity with a tier 1 AI platform provider is expected to enter volume production in 2027. Second, near packaged optics, or NPO, chipsets enabling multi-rack scale-up clusters via optical links are also targeted for 2027 production. Third, fully integrated Scorpio X-Series fabric switches with CPO optical engine modules represent our longer-term roadmap for 2028 and beyond, supporting larger fabric domains with terabits per second optical I/O per switch. Beyond standard products, we continue to make solid progress with our custom solutions. The increasingly diverse XPU landscape and workloads are driving demand for custom connectivity solutions. We are executing across multiple vectors with custom solutions for

hybrid rack architecture, including NVLink fusion-based designs and, separately, custom Leo memory controllers in specialized interface applications. We expect initial shipments for these custom design wins to commence in 2027. We are also engaged with multiple customers on additional high-value custom connectivity opportunities to deliver end-to-end purpose-built silicon solutions. Together, these design wins and active opportunities validate our capability to deliver differentiated custom silicon and represents a multi-billion dollar incremental opportunity over time. Overall, our platform philosophy, which we call AI Your Way, reflects our strategy. Rather than prescribing a single architecture approach, we provide the broadest portfolio of connectivity solutions that let customers deploy AI infrastructure optimized for their specific workloads, whether that's standards-based or custom or copper or optical interconnects and scale-up or scale-out connectivity across a plethora of standards. This architectural flexibility is becoming a competitive advantage as the market diversifies. To summarize, Astera Labs has hit another important inflection point with Scorpio high-radix AI fabrics reaching volume production. This marks our evolution to a complete AI fabric infrastructure provider, extending beyond signal conditioning and memory connectivity solutions. Our expansion into optical interconnects and custom solutions further broadens our addressable market in 2027 and beyond. With that, let me turn over the call to Des to walk through our Q2 financial results and Q3 guidance in more detail.

Desmond Lynch: Thank you, Sanjay, and good afternoon, everyone. Today, I will review our Q2 financial results and then discuss our Q3 guidance, both presented on a non-GAAP basis. We delivered outstanding second quarter results with record revenue of \$392.4 million, up 27% sequentially and up 104% year-over-year. Growth was broad-based across our AI fabric and signal conditioning portfolios, led by PCIe 6.0 products and the initial production ramp of Scorpio X-Series. Our Scorpio product family delivered a strong Q2. Scorpio P-Series momentum continued during the quarter, driven by the scaling of existing customer deployments and the early ramp of new programs into production. Scorpio X-Series also grew strongly as we began shipping initial production volumes across a variety of radix configurations. Aries revenue reached a new quarterly record driven by strong PCIe 6.0 retimer adoption across both scale-up and scale-out applications. Overall, revenues from PCIe 6.0 across our AI fabric and signal conditioning portfolios represented more than 50% of total company revenue in Q2, up from one-third in Q1. Taurus's revenue also grew strongly in Q2, driven by increased unit shipments across both AI and general-purpose compute platforms. Non-GAAP gross margins for the second quarter were 73.7%, above our guidance of 73%. Non-GAAP operating expenses for the second quarter were \$135.8 million, up 10% sequentially, reflecting continued investment in our roadmap. Non-GAAP operating margins for the second quarter were 39.1%, up 290 basis points from Q1, driven primarily by strong revenue growth and the resulting operating leverage. For the second quarter, interest in other income was \$12.1 million. Our non-GAAP tax rate was 12%, and non-GAAP net income was \$145.8 million, resulting in non-GAAP diluted earnings per share of \$0.80, an increase of more than 30% from Q1. We enter the quarter with cash, cash equivalents to marketable securities, totaling \$1.25 billion, of \$68.5 million from Q1 driven by cash of \$87.7 million. Turning to our outlook for the third quarter. We expect revenue to be in the range of \$540 million-\$560 million. At the midpoint, this represents 40% sequential growth driven by the transition of Scorpio X-Series to volume production, continued strength in Aries PCIe 6.0 retimers, and pre-production shipments of Taurus solutions supporting 100 gigabits per lane for 800-gigabit Ethernet. As Jitendra mentioned, we expect Scorpio to become our largest product line by revenue in the third quarter, marking an important milestone in our journey from smart signal conditioning products to mission-critical AI fabric solutions. This milestone reflects the progress we have made over the past several quarters as Scorpio X-Series transitions to volume production. We expect third quarter non-GAAP gross margin to be approximately 72%. We expect third-quarter non-GAAP operating expenses to be between \$156 million and \$160 million. Our non-GAAP operating margin is expected to be approximately 43%, up 400 basis points sequentially, which reflects significant operating leverage as our revenue expands. Interest in other income is expected to be approximately \$12 million, and we expect our non-GAAP tax rate to be approximately 12%. We expect our Q3 diluted share count to be approximately 185 million shares outstanding. Overall, we expect non-GAAP fully diluted earnings per share to be between \$1.16 and \$1.21. In closing, Q2 was an excellent quarter with strong execution, driving record revenue and expanding operating leverage. Our strong Q3 guidance reflects continued momentum and confidence in sustained growth across both our AI fabric and signal conditioning portfolios. We will continue to invest strategically to capture the substantial opportunity ahead while maintaining strong profitability and driving long-term value to our stockholders. This concludes our prepared remarks. I will now turn the call back to our operator to begin Q&A. Operator?

Operator: We will now begin the question-and-answer session. Please limit yourself to one question. At this time, I would like to remind everyone, in order to ask a question, press star then one on your telephone keypad. We'll

pause for just a moment to compile the Q&A roster. Your first question comes from the line of Harlan Sur with J.P. Morgan. Harlan, your line is now open. Please go ahead.

Harlan Sur: Good afternoon. Thanks for taking my question and great execution by the team. It's great to see the volume production ramp of your Scorpio X scale-up fabric switching solutions here in the third quarter. I think in the press release you said that Scorpio X ramp would be led by your newer 320-lane products. I think that's with one hyperscaler customer, right? Do you have another hyperscale cloud customer that is targeted to ramp your lower-rate X, Scorpio X product family, as well? Is that more of a Q4 high-volume ramp? With the strong ramp in Q3, is Scorpio X revenues going to exceed Scorpio P revenues this quarter, or is that still looking to be more like the Q4 timeframe? Thank you.

Sanjay Gajendra: Thanks, Harlan. A couple of points. First, the Scorpio X-Series has been gaining tremendous amount of interest in the market. We have started shipping the parts for pre-production for multiple customers starting this quarter. We do expect Q3 at least for our lead opportunity and lead customer to ramp to high-volume production in Q3 and essentially taking over, or surpassing the revenue from Scorpio P-Series. Which should give you an idea of the value that we bring in scale-up use cases. We do expect to ship to revenue for additional customers on Scorpio X by end of the year. In general, what I would say is that there's been an exceptional traction that we're seeing both for Scorpio X and P. X tends to be more tied to scale-up and various different configurations. P goes to a multitude of opportunities. Overall it's been great to see how the customer base is seeing the value that we deliver with our Scorpio family. Like we noted in the call, Q3 onwards, we expect Scorpio as a family to be the largest product line within the company.

Harlan Sur: Thank you, Sanjay.

Operator: Your next question comes from the line of Blayne Curtis with Jefferies. Blayne, your line is now open. Please go ahead.

Blayne Curtis: Hey, guys. Thanks for taking my question and great results. I know there's going to be a lot on Scorpio, I want to ask on the Aries product; you said record revenue. I think when you gave initial outlook, it was directionally growth. I think people had pretty low expectations. Just kind of curious if you could talk about the drivers within Aries, and then as you look to kind of the next generations of AI servers, what kind of design traction are you seeing and retimer attach in those systems?

Sanjay Gajendra: Yeah. Two parts to that. One is, again, Aries 5, which is our gen 5 portfolio that's been around; I want to say definitely continues to expand, especially as we're seeing inference-based applications start showing up. That's driving a new wave of growth, is how I would put it. Aries 6, which is for gen 6, supporting the higher data rate, the 64 gigabits per second. Obviously that comes in with higher attach rate and going into several new platforms. Also, there are other use cases around the cabled application and so on. Overall, like we've always said, the retimer is a socket that somewhat keeps on giving in the sense that you need it for a variety of use cases. If anything, the surprise is how long Gen 5 has continued, and given that there is a new set of use cases opening up for inference, all of those are contributing to help us achieve a new record on the Aries portfolio.

Jitendra Mohan: By the way, we expect this trend to continue as new Gen 6 platforms and GPUs are being developed. This is not a trend that will stop. Then, of course, then we move on to Gen 7. You can expect some of this cycle to continue for a long time.

Blayne Curtis: Thank you.

Operator: Your next question comes from the line of Joseph Moore with Morgan Stanley. Joe, your line is now open. Please go ahead.

Joe Moore: Great. Thank you. You mentioned UALink in your roadmap there. Can you talk a little bit about that roadmap and, do you see, I guess, a transition from PCI to UALink and how broad do you expect the customer adoption to be there?

Jitendra Mohan: Hey, Joe. Thank you for the question. Yeah, we continue to see a lot of interest in UALink for scale-up networking. As you mentioned during the prepared remarks, the ecosystem is evolving. We have new IP vendors coming online, new tools coming online. Some of the other vendors have also introduced their own products or announced their own products. That kind of all points to a nice, vibrant ecosystem. From a customer standpoint, we continue to stay engaged with our customers. There is continuing interest in UALink. A lot of it is actually driven by not only the additional throughput that UALink brings but also because of the low

latency that UALink will deliver compared to some of these other protocols. We find a lot of traction for that in inferencing applications. The throughput and the latency and the open ecosystem is definitely continuing to drive the traction from our customers. We are completely plugged in terms of the development cycle for UALink. We plan to intercept the XPU's that are on track to be released in 2027 with our own UALink solutions. As we have discussed, we want to come in here with a full portfolio of devices, which would be anchored around our switch platform. Then also providing signal conditioning solutions both for copper and optical. In fact, we recently introduced our Taurus family, which supports UALink at 200 gig per lane line rates. It's all looking very good. We continue to stay completely engaged, as things move from the PCI Express generation to UALink, we do expect the dollar content per XPU to increase; that's going to be a favorable trend for us.

Joe Moore: Great. Thank you.

Operator: Your next question comes from the line of Tore Svanberg with Stifel. Your line is now open. Please go ahead.

Tore Svanberg: Yes. Thank you; congratulations on the strong results. You talked about new software for the Scorpio X deployment. I think you even talked about really improving efficiency, tokenomics, and so on and so forth. Can you maybe elaborate a little bit on that? I'm especially curious if that new software is also landing you more new customers, because I assume that's something that's going to be very beneficial as they launch their switch fabrics.

Jitendra Mohan: Yeah. As Sanjay mentioned, Tore, there is exceptional traction from our new customers for the new software devices that we introduced. Part of it comes just because of the fact that we have higher radix and higher lane count; therefore, we can connect more XPU's to the switch. A lot of it is also stemming from the advanced capabilities that we built in, I think this is what you refer to when you talk about software, which is the in-network compute as well as HyperCast. Both of these are very important technologies that really set us apart and enable very low latency inferencing as well as very high efficiency for in-network compute. Almost 2x improvement in the collective operations that are required for training. All of these are hardware-based features, but they are enabled through software, and this software resides in our COSMOS platform. While previously COSMOS was mainly responsible for optimization, telemetry, diagnostics, now we are becoming part of the orchestration layer. The XPU's are talking directly to our COSMOS platform to run the workloads. This drives a lot more stickiness into our solution. Many customers that have adopted our platform in this generation will continue to use this for the next generation. It's the easiest thing for them to do. We continue to support these types of advanced capabilities in our current devices, and we'll carry them forward for the next generation devices as well.

Tore Svanberg: That's very clear. Thank you.

Operator: Your next question comes from the line of Natalia Winkler with UBS. Natalia, your line is open. Please go ahead.

Natalia Winkler: Hi. Thank you very much for taking my question. I wanted to follow up on the CXL opportunity. You guys mentioned the CXL controller and kind of two new hyperscalers that you will be ramping in 2027. Curious how that CXL opportunity plays out for the switching content as well, and how do you see that market kind of progressing from here?

Sanjay Gajendra: Yeah, I think it was about time, isn't it? I think it's finally great to see some traction back on CXL. As you know, we've been investing in the technology for quite a while now, and it's great in some ways because we understand exactly what is needed both from a hardware and software standpoint. Now, we see actually three main use cases. One is largely driven around the memory market dynamics. People wanting to find alternates to the availability or the high cost of memory today. They're trying to do that through CXL and try to leverage the fact that you can stick different kinds of memory behind a CXL bus. The second one, which is very exciting for us, is the GPU attach or accelerator attach, where combined with our KV cache acceleration type of function, we're able to do unique things that fundamentally are improving the latency and performance for inference applications. The third one is something that we've been talking for a while, which is general compute or HPC application, where for use cases like SAP, where you need a high memory or memory-intensive workload, the CXL provides a way of adding more memory to the CPU. Those are the three main applications or dynamics that we see. For us, what we are tracking is, of course, multiple design win and opportunities at this point, with most of them in qualification or early stage this year and getting into high-volume production in 2027. We do see

this trend to continue, by the way, and we are investing heavily in coming out with technologies and capabilities to solve the memory wall problem and the memory availability problem in some ways, given where the market is on that front. Overall, we are positive and we are continuing to invest.

Natalia Winkler: Thank you.

Operator: Your next question comes from the line of Sean O'Loughlin with TD Cowen. Sean, your line is now open. Please go ahead. Sean O'Loughlin: Hey, guys. Congrats on another great set of results, and thanks for letting me hop on. I wanted to ask maybe a bigger picture question on the product portfolio, especially as we start to think about optics really starting to be something that sees at-scale deployment within most investors' time horizons. I guess the question is, as I think about a UALink, NPO, or CPO switch, you have an optics-based Ethernet PHY. Is there a fundamental reason why Taurus over time couldn't also transition that physical media-dependent layer to one that's based on fiber rather than copper? Does that, in effect, extend the assumed lifetime revenue of what maybe most of us are thinking of as an inherently copper-based solution today? Apologies if I'm leading the witness here a bit, but I just wanted to get your thoughts on that.

Sanjay Gajendra: No, that's a great question. Let me comment on that. First of all, we do fully expect copper and optical to coexist. As you know, most of the scale-up connectivity today is done in copper, and our customers are asking us to continue to push that forward, and we will continue to do that. As these data rates go up and the signal integrity challenges increase, there is definitely room for additional products and additional product categories, just like we announced our Taurus 200G, both retimers as drivers in support of these applications. We do expect that for this generation and the next generation, the scale-up within a rack will continue to be copper. As these cluster sizes grow and the data rates increase, optical will definitely play an increasingly important role as you start connecting these racks to each other and cluster size expands from maybe one or two racks to four, five, or even a larger number of racks. We are very well positioned to play our part in this developing market segment. We are very closely engaged with our customers. We understand when they want to deploy these increased cluster sizes, and we expect that the first deployment will happen with NPO solution. As we go from copper, in addition to copper, to NPO solution, our dollar content per XPU or per link actually goes up. We definitely welcome this change. Overall, if you look at the optical market, it's several tens of billions of dollars in TAM that gets unlocked. Once NPO gets deployed and data rates continues to increase, we will start to see CPO getting deployed. We look at it as 2027 being the year where NPO gets deployed; 2028 and beyond is when CPO gets deployed. From a product development standpoint, we are very comfortable. We started investing in this a couple of years back. We have a team of engineers working on the EIC solution, which is the electrical component. We have a team of engineers that is working on the silicon photonics. Of course, late last year, we acquired aiXscale Photonics, which gives us the ability to develop a connector. Put all of these things together, we have now the ability, or we will have the ability, to have full optical engines, which will eventually lead to not only optical links and Taurus-like components, as you pointed out, but actually a fully optically enabled Scorpio switch. We are really looking forward to that, supplying both an optically enabled Scorpio switch as well as the individual components that make up the full optical link. That's a tremendous opportunity for us in addition to what we have at copper. Sean O'Loughlin: Thanks for that. Congrats again.

Sanjay Gajendra: Thank you.

Operator: Your next question comes from the line of Papa Sylla with Citigroup. Papa, your line is now open. Please go ahead.

Papa Sylla: Thank you for taking my question. Congrats on the impressive results. I wanted to double-click on the content for accelerator. I think when you initially kind of IPO'd, it was more around the \$50-\$100 for accelerator, and over the past quarters or so, it went to \$1,000+. Sanjay, you mentioned today Scorpio X alone could be well over \$1,000 per accelerator. I'm wondering if that \$1,000+ for Scorpio X is it a pre-UALink metric. Perhaps if you can provide color on the path to that \$1,000+ per accelerator for Scorpio X. Finally, maybe tying everything together, if you can speak more generally on your content per accelerator currently, and where do you see that growing over the next 2-3 years?

Sanjay Gajendra: Yeah, a good question. I think this is a fundamental question that you asked, and it's a very important focus for us, which is how do we keep ensuring that our long-term business continues to grow, and we are able to sustain the kind of growth that we're able to deliver? It really, like you highlighted, comes down to how can we keep increasing the dollar content that we can get per XPU attached. Like you correctly noted, we started with small numbers when IPO time. Now we see our business contributing multiple thousands of

dollars per XPU with the X scale or scale-up switch itself offering over \$1,000 per XPU. This is already starting to happen, of course, with the high-radix switch that we are shipping right now. At the same time, this is just the beginning of it, because when you start adding several other features that are important in terms of optical, for example, or adding more signal conditioning because of the higher speed and the higher attach rate, we do believe that this number will continue to scale up. That is what we are focusing our investment right now, which is how do we ensure, A, we keep increasing the dollar content we get per XPU targeting multiple thousands of dollars, and B, how do we continue to keep diversifying our business, adding more hyperscalers around the anchor socket that we have, which is a scale-up switch, and continue to grow significantly above where the market is.

Papa Sylla: Got it. Very helpful. Thank you.

Operator: Your next question comes from the line of Tom O'Malley with Barclays. Tom, your line is now open. Please go ahead.

Trip Smith: Hi, team. This is Trip Smith on for Tom O'Malley. Appreciate the question and very nice results. I was just curious about how we should think about Scorpio's gross margins now that it's becoming a bigger part of the business and what are the puts and takes there for gross margins looking out. I believe you still have the Amazon warrant impact, so I just wanted to double-check on that and how we should be modeling that over the next few quarters. Thank you.

Desmond Lynch: Hey, thanks for the question. It's Des here. As we mentioned in the prepared remarks, we do expect our Q3 gross margins to be approximately 72%. Really, as our portfolio continues to sort of diversify, we continue to see a wider range of gross margins, with the main impacts being silicon versus module revenue mix and the different use cases associated with our switching portfolio. On Scorpio across both P and X Series, we do have a wide range of margins, really driven by use case and lane count utilization. I would say, on average, Scorpio across all use cases has an average margin profile broadly in range with our corporate gross margins. In any given quarter, this mix can change. What I would say on the switching side is that we are going after large opportunities, and as the top line continues to scale, we'll see the gross margin dollars flow through and contribute to the bottom line. We clearly saw that in our Q2 results and also our Q3 guidance from here. Really, given the broadening of the portfolio, I would expect our gross margins continue to trend towards a long-term target of 70% from here.

Operator: Your next question comes from the line of Ananda Baruah with Loop Capital Markets. Ananda, your line is now open. Please go ahead.

Ananda Baruah: Thanks, operator, for the question. Good morning, I mean, good afternoon, guys. Thanks for taking the question. What's a useful way to think about the contribution that the China hyperscale can make across the portfolio as you continue to introduce the new products and move up from generation to generation? Thanks.

Sanjay Gajendra: Ananda, it's a good observation. In China, as you can imagine, open protocols are quite popular, which includes PCI Express, Ethernet, and increasingly, NVLink. We definitely see that as a good market, not nearly as big as what we have here, but certainly there is a lot of interest, and we are engaging with the customers there with our PCI solutions as well as our Taurus-based Ethernet solution. In aggregate, that can be a very promising opportunity, especially as PCI Express is the nervous system of servers and probably the best way to connect many add-in card-format GPUs together to deliver good inferencing capabilities.

Ananda Baruah: I appreciate. Is it a way that can be noticeable in the PNL over time?

Sanjay Gajendra: I think it'll be in an absolute scale; yes, our revenues from China will continue to increase. However, we do expect the rest of the world to increase even faster. As a percentage of total revenue, they will continue to be in the single digits.

Ananda Baruah: Got it.

Operator: Your next question comes from Suji Desilva with Roth Capital Partners. Suji, your line is now open. Please go ahead.

Suji Desilva: Hi, Jitendra. Hi, Jitendra, Sanjay, and Des. Congrats on the results here. I'm curious; you talked about the software supporting telemetry and now orchestration gaining more kind of functionality there. Is there a kind of a soft line where you might cross over to helping the XPU itself with some of the functionality as a

coprocessor, or would you draw sort of a straight line there as connectivity versus that eventual trend in memory, compute, and things like that?

Sanjay Gajendra: Very good point. Actually, that is what I was trying to say when I answered the question earlier: that with the advancement of the in-network compute and Hypercast features in our latest Scorpio 320-lane device, we are indeed crossing over that line where the XPU is directly talking to the switch to orchestrate the workloads, to reduce the latency, and increase the throughput for both inferencing as well as training workloads. That's a very important development for our COSMOS platform, where not only are we doing the traditional optimization, customization, diagnostic, and telemetry, but we are really helping to improve the workload and directly improve the utilization of the GPUs that are connected to our switch. It's a very important development, and as I mentioned, it makes our software solution more powerful and a lot more sticky as customers designing our current generation of switches and then look to upgrade them in the future.

Suji Desilva: Okay, thanks.

Operator: Your next question comes from the line of Simon Leopold with Raymond James. Simon, your line is now open. Please go ahead.

Jeff Koche: Yeah, thanks. Jeff Koche in for Simon. Maybe you could just talk qualitatively about your end market concentration. How has it trended the quarter? How do you expect it to trend over the next, I guess, through 2027, especially with the new launches?

Desmond Lynch: Yeah, it's a really good question. What I was really pleased to see in our first half is that, really, it's been a broad-based strength across our product portfolio. Really looking at our drivers, I would say Aries delivered the record quarterly revenue in Q2 and remained our largest product line, really supported by the ongoing transition to PCIe 6.0 across both scale-out and scale-up connectivity. I do expect Aries to continue to grow in Q3 and beyond. As we mentioned in the prepared remarks, Scorpio continues to show strong momentum in Q2, really driven by X and P Series. As we mentioned, we do expect to see Scorpio becoming our largest product line in terms of revenue as a result of the production ramp of the high-radix solutions in Q3. Taurus also continues to show nice growth, really supported by AI and general-purpose platforms. I do expect to see further growth in the third quarter, driven by the strong base of the 400 gig demand, as well as the early contributions of 800 gigs as well. I think as Sanjay mentioned, we continue to see strong progress on the Leo CXL products, and we continue to see some revenue here in the back half of the year, which will ramp more materially into 2027. As we look into 2027, we'll continue to see further opportunities on the optical and custom side. We're really pleased with the overall performance of the portfolio and the continued diversification.

Jeff Koche: Great.

Operator: Our next question comes from Karl Ackerman with BNP Paribas. Karl, your line is now open. Please go ahead.

Karl Ackerman: Yes, thank you. Jitendra, Scorpio X is ramping significantly in the second half of this year, largely driven by the product transition at your main customer. Could you discuss the breadth of customer wins, including a growing mix of neo clouds, and whether your Scorpio X opportunity extends across both liquid-cooled and air-cooled servers? Thank you.

Sanjay Gajendra: Yeah, I'll take this question. It's Sanjay here. Like we've said, the demand for Scorpio X and Scorpio P has truly been exceptional. We have over 10 customers engaged on the Scorpio X itself. Many of these opportunities are either getting into pre-production or getting deeper into the qualification cycle. To that standpoint, we are already shipping Scorpio X to multiple customers right now as they get ready for their own internal qualification cycle. We do expect that we'll have additional customers ramping to production towards the end of this year, early next year, based on the progress that they can make. This trend is something that, of course, given the features and capabilities that we have, especially with the 320-lane, we're able to unlock several things when it comes to inference-type applications. There are several unique things that we are able to achieve, and all of these are contributing to diversifying our customer base on Scorpio X. Like I noted on Scorpio P, it's truly broad-based because the P series go into standard PCI switching applications. From storage to networking and some of the scale-out use cases, we have several design wins and design-ins that are ramping to production. In general, I want to say Scorpio family is taking off in a very positive way for us, and we intend to add more customers and more revenue growth in the second half through 2027.

Jitendra Mohan: If I may just add to that, the chip itself supports both liquid cooling and air cooling. In fact, we do have customers that are designing it for both of them. Your point about the neo clouds is also very valid, because people, when they start to deploy these solutions for especially inferencing applications, sometimes they can do that without building a full rack-level solution. We do see this increasing use case of enterprise customers and neo-clouds looking at solutions that are based on Scorpio as a scale-up and using add-in card form factor GPUs to deliver good inferencing TCO.

Karl Ackerman: Thank you.

Sanjay Gajendra: Also maybe to add one more point that is somewhat lost is even with the lead hyperscaler we have, there are multiple platforms where we are designed in. It's not like there is just one, right? I think that's a nuance that needs to be considered as you look at our business.

Operator: There are no further questions at this time. I will now turn the call back over to Leslie Green for closing remarks.

Leslie Green: Thank you, everyone, for your participation and questions. Please refer to our investor relations website for information regarding upcoming financial conferences and events. Thanks so much.

Operator: This concludes today's conference call. You may now disconnect. Read Next Aug 11, 2026 By John Ballard Astera Labs vs. Arm: Which Technology Stock Is a Better Buy in 2026? Aug 6, 2026 By Sara Appino Astera Labs vs. AppLovin: Which Technology Stock Is a Better Buy in 2026? Aug 6, 2026 By Trevor Jennewine Nvidia, Micron, and Sandisk Are Up Over 1,000% Since the AI Boom Started. These 2 AI Stocks Could Be the Next Big Winners. Aug 1, 2026 By John Ballard Astera Labs vs.

Arista Networks: Which Tech Stock Is a Better Buy in 2026? Jul 31, 2026 By Scott Levine Why Astera Labs Stock Rocketed Higher Today Jul 31, 2026 By John Ballard Astera Labs vs. Intel: Which Technology Stock Is a Better Buy in 2026? This article is a transcript of this conference call produced for The Motley Fool. While we strive for our Foolish Best, there may be errors, omissions, or inaccuracies in this transcript. Parts of this article were created using Large Language Models (LLMs) based on The Motley Fool's insights and investing approach. It has been reviewed by our AI quality control systems. Since LLMs cannot (currently) own stocks, it has no positions in any of the stocks mentioned. As with all our articles, The Motley Fool does not assume any responsibility for your use of this content, and we strongly encourage you to do your own research, including listening to the call yourself and reading the company's SEC filings. Please see our Terms and Conditions for additional details, including our Obligatory Capitalized Disclaimers of Liability. The Motley Fool recommends Astera Labs.